

EVM (part 2)

Recap: Ethereum consensus

consensus layer (beacon chain)

Today: EVM

compute layer (execution chain): The EVM

consensus layer (beacon chain)

Ethereum compute layer: the EVM

World state: set of accounts identified by 32-byte address.

Two types of accounts:

(1) externally owned accounts (EOA):

controlled by ECDSA signing key pair (pk,sk).

sk: signing key known only to account owner

(2) contracts: controlled by code.

code set at account creation time, does not change

Data associated with an account

| <u>Account data</u> | <u>Owned (EOA)</u> | <u>Contracts</u> |
|------------------------------|--------------------|----------------------------------|
| address (computed): | $H(pk)$ | $H(CreatorAddr, CreatorNonce)$ |
| code : | \perp | CodeHash |
| storage root (state): | \perp | StorageRoot |
| balance (in Wei): | balance | balance (1 Wei = 10^{-18} ETH) |
| nonce : | nonce | nonce |

 (#Tx sent) + (#accounts created): anti-replay mechanism

State transitions: Tx and messages

Transactions: signed data by initiator

- **To:** 32-byte address of target (0 → create new account)
- **From, [Signature]:** initiator address and signature on Tx (if owned)
- **Value:** # Wei being sent with Tx (1 Wei = 10^{-18} ETH)
- Tx fees (EIP 1559): **gasLimit, maxFee, maxPriorityFee** (later)
- if To = 0: create new contract **code = (init, body)**
- if To ≠ 0: **data** (what function to call & arguments)
- **nonce:** must match current nonce of sender (prevents Tx replay)
- **chain_id:** ensures Tx can only be submitted to the intended chain

State transitions: Tx and messages

Transaction types:

owned → owned: transfer ETH between users

owned → contract: call contract with ETH & data

| Transaction Hash | Method | Block | Age | From | To | Amount | Txn Fee |
|----------------------------------|--------------------|--------------------------|-------------|--|---|-----------------|------------|
| 0x427c31b0a2... | Transfer | 22334348 | 15 secs ago | beaverbuild | 0x123B9747...C68A0A806 | 0.010695384 ETH | 0.00000883 |
| 0xe208fcc8c90... | Transfer | 22334348 | 15 secs ago | 0x9Bd017bc...3f27E661e | 0x571AD708...73fa8766C | 0.00010801 ETH | 0.00000885 |
| 0x1def593585c... | Transfer | 22334348 | 15 secs ago | 0x906457C3...D2f990e1E | 0xEf5599e...28121eD24 | 0.04994585 ETH | 0.00000885 |
| 0x2e25b793a1... | Transfer | 22334348 | 15 secs ago | 0x016606Ac...44B4049Ee | 0xC1BEd8E8...835EF8B09 | 0.182092191 ETH | 0.00000886 |
| 0x838bd63745... | Pre Collaterali... | 22334348 | 15 secs ago | 0x57Dd832f...D8818bF4d | 0x3A4F41da...71bF4Cc59 | 0 ETH | 0.00003483 |
| 0xf60eb42b563... | Transfer | 22334348 | 15 secs ago | 0x68B0F72D...61E84f7f5 | 0x3CAcB76c...C6f71cFB1 | 0.009995342 ETH | 0.00000898 |
| 0x1cea36b658... | Transfer | 22334348 | 15 secs ago | 0xf87BDC2b...0Dde8AF6b | 0x3CAcB76c...C6f71cFB1 | 0.009988976 ETH | 0.00000898 |
| 0xc9c9b3b7aa... | Transfer | 22334348 | 15 secs ago | 0x1242f107bf3412a893b5f47b9b1a7b539e37822f | 0x3CAcB76c...C6f71cFB1 | 0.052833105 ETH | 0.00000898 |
| 0x8b1a4a52d8... | Swap | 22334348 | 15 secs ago | 0xa84B9E91...b4187Cc08 | Aggregation Router V6 | 0.03 ETH | 0.00008007 |
| 0x20ba850ab8... | Transfer | 22334348 | 15 secs ago | 0xEDCa7976...A8C8503c1 | 0xf066CEd2...A3C9b977F | 0.024854505 ETH | 0.00000903 |
| 0x3cece221e2... | Transfer | 22334348 | 15 secs ago | 0x3CAcB76c...C6f71cFB1 | 0xDdEc41b0...0a08fE03c | 0.01 ETH | 0.00000904 |
| 0x8eaa3695daf... | Transfer | 22334348 | 15 secs ago | Cobo: Custody | 0x0D65fEE7...F4822b078 | 0.4495 ETH | 0.00000905 |
| 0xbd1d8e3b8df... | Transfer | 22334348 | 15 secs ago | 0x97Ca787E...0Ec2DF5b9 | Tether: USDT Stablecoin | 0 ETH | 0.00001767 |
| 0x256605a075... | Transfer | 22334348 | 15 secs ago | 0x681fF4BF...8C9017Bd9 | Tether: USDT Stablecoin | 0 ETH | 0.00001775 |
| 0xf3b21d58ad7... | Transfer | 22334348 | 15 secs ago | 0x9320f4d9...B153B82a5 | Circle: USDC Token | 0 ETH | 0.00001945 |
| 0x7cc7874007... | 0x59805e69 | 22334348 | 15 secs ago | 0xE0b7DEab...41D058C79 | 0xE54E03fC...0D8CB2dAf | 0 ETH | 0.00002283 |

Messages: virtual Tx initiated by a contract

Same as Tx, but no signature (contract has no signing key)

contract → EOA: contract sends funds to user

contract → contract: one program calls another (and sends funds)

One Tx from user: can lead to many Tx processed. Composability!

Tx from owned addr → contract → another contract

Tx from owned addr → contract → another contract → EOA

Example Tx

| State | |
|---|-----------------|
| 14c5f8ba: - 1024 eth | <u>owned</u> |
| bb75a980: - 5202 eth if !contract.storage[tx.data[0]]: contract.storage[tx.data[0]] = tx.data[1] [0, 235235, 0, ALICE | <u>contract</u> |
| 892bf92f: - 0 eth send(tx.value / 3, contract.storage[0]) send(tx.value / 3, contract.storage[1]) send(tx.value / 3, contract.storage[2]) [ALICE, BOB, CHARLIE] | <u>contract</u> |
| 4096ad65: - 77 eth | <u>owned</u> |

world state (four accounts)

updated world state

Example Tx

| State | |
|---|-----------------|
| 14c5f8ba: - 1024 eth | <u>owned</u> |
| bb75a980: - 5202 eth if !contract.storage[tx.data[0]]: contract.storage[tx.data[0]] = tx.data[1] [0, 235235, 0, ALICE | <u>contract</u> |
| 892bf92f: - 0 eth send(tx.value / 3, contract.storage[0]) send(tx.value / 3, contract.storage[1]) send(tx.value / 3, contract.storage[2]) [ALICE, BOB, CHARLIE] | <u>contract</u> |
| 4096ad65: - 77 eth | <u>owned</u> |



Transaction

From:
14c5f8ba

To:
bb75a980

Value:
10 eth

Data:
2,
CHARLIE

Sig:
30452fdedb3d
f7959f2ceb8a1



| State' | |
|---|--|
| 14c5f8ba: - 1014 eth | |
| bb75a980: <u>- 5212 eth</u> if !contract.storage[tx.data[0]]: contract.storage[tx.data[0]] = tx.data[1] [0, 235235, <u>CHARLIE</u> , ALICE .. | |
| 892bf92f: - 0 eth send(tx.value / 3, contract.storage[0]) send(tx.value / 3, contract.storage[1]) send(tx.value / 3, contract.storage[2]) [ALICE, BOB, CHARLIE] | |
| 4096ad65: - 77 eth | |

world state (four accounts)

updated world state

An Ethereum Block

Block proposer creates a block of n Tx: (from Txs submitted by users)

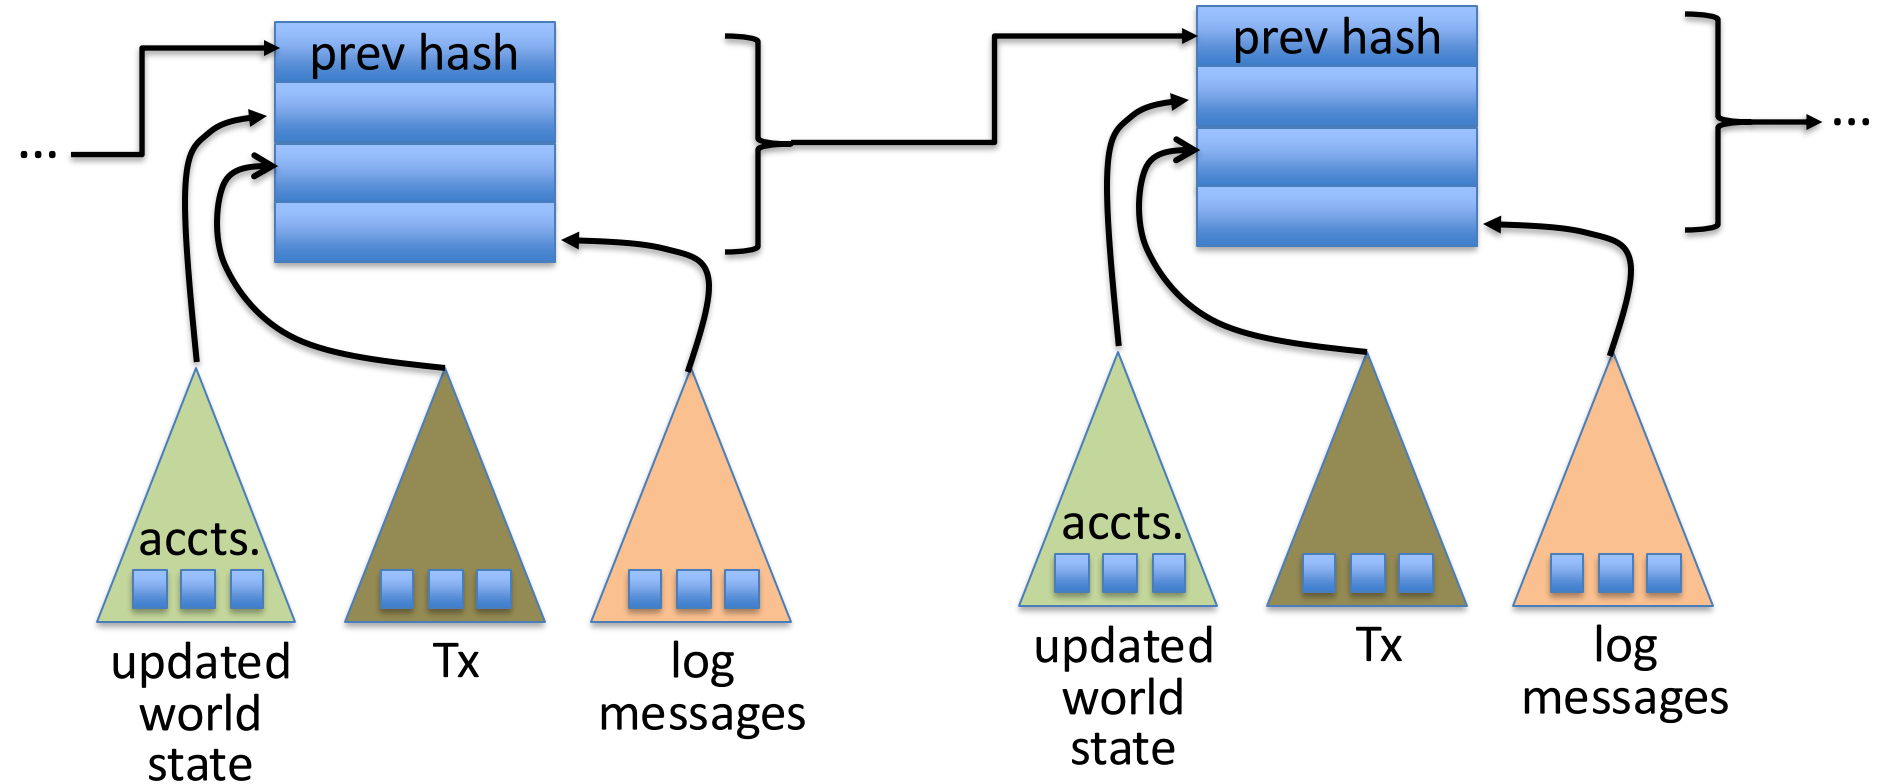
- To produce a block do:
 - for $i=1,\dots,n$: execute state change of Tx_i sequentially
(can change state of $>n$ accounts)
 - record updated world state in block

Other validators re-execute all Tx to verify block \Rightarrow
sign block if valid \Rightarrow enough sigs, epoch is finalized.


Block header data (simplified)

- (1) consensus data: proposer ID, parent hash, votes, etc.
- (2) address of gas beneficiary: where Tx fees will go
- (3) world state root:** Merkle hash of all accounts in the system
- (4) Tx root:** Merkle hash of all Tx processed in block
- (5) Tx receipt root:** Merkle hash of log messages generated in block
- (5) Gas used: used to adjust gas price (target 15M gas per block)

The Ethereum blockchain: abstractly



Amount of memory to run a node




An example contract: NameCoin

```
contract nameCoin {  
  
    struct nameEntry {  
        address owner; // address of domain owner  
        bytes32 value; // IP address  
    }  
  
    // array of all registered domains  
    mapping (bytes32 => nameEntry) data;  
  
    // event emitted when registering domain  
    event Register(address indexed owner, bytes32 name);
```


An example contract: NameCoin

```
function nameNew(bytes32 name) {  
    // registration costs is 100 Wei  
  
    if (data[name] == 0 && msg.value >= 100) {  
        data[name].owner = msg.sender    // record domain owner  
        emit Register(msg.sender, name)  // log event  
    }  
}
```



Code ensures that no one can take over a registered name

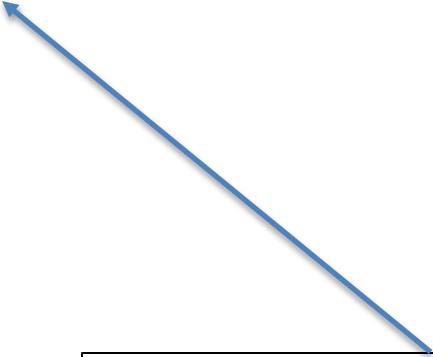
Serious bug in this code! Front running. Solved using commitments.

An example contract: NameCoin

```
function nameUpdate(bytes32 name, bytes32 newValue, address newOwner) {  
    // check if message is from domain owner,  
    //                and update cost of 10 Wei is paid  
    if (data[name].owner == msg.sender    &&    msg.value >= 10) {  
        data[name].value = newValue;        // record new value  
        data[name].owner = newOwner;        // record new owner  
    }  
}
```

An example contract: NameCoin

```
function nameLookup(bytes32 name) {  
    return data[name];  
}  
  
} // end of contract
```



Used by other contracts
Humans do not need this
(use etherscan.io)

EVM mechanics: execution environment

Write code in Solidity (or another front-end language)

⇒ compile to EVM bytecode

(some projects use WASM or BPF bytecode)

⇒ validators use the EVM to execute contract bytecode
in response to a Tx

The EVM

Stack machine (like Bitcoin) but with JUMP

- max stack depth = 1024
- program aborts if stack size exceeded; block proposer keeps gas
- contract can create or call another contract

In addition: two types of zero initialized memory

- Persistent storage (on blockchain): SLOAD, SSTORE (expensive)
- Volatile memory (for single Tx): MLOAD, MSTORE (cheap)
- LOG0(data): write data to log

see <https://www.evm.codes>

Every instruction costs gas, examples:

SSTORE **addr** (32 bytes), **value** (32 bytes)

- zero → non-zero: 20,000 gas
- non-zero → non-zero: 5,000 gas (for a cold slot)
- non-zero → zero: 15,000 gas refund (example)

Refund is given for reducing size of blockchain state

CREATE : $32,000 + 200 \times (\text{code size})$ gas;

CALL **gas**, **addr**, **value**, **args**

SELFDESTRUCT **addr**: kill current contract (5000 gas)

Gas calculation

Why charge gas?


- Tx fees (gas) prevents submitting Tx that runs for many steps.
- During high load: block proposer chooses Tx from mempool that maximize its income.

Old EVM: (prior to EIP1559)


- Every Tx contains a gasPrice ``bid`` (gas \rightarrow Wei conversion price)
- Producer chooses Tx with highest gasPrice ($\max \sum(\text{gasPrice} \times \text{gasLimit})$)
 \Rightarrow not an efficient auction mechanism (first price auction)

Gas prices spike during congestion


GasPrice in Gwei:



Average Tx fee in USD



Gas prices spike during congestion



Gas calculation: EIP1559

EIP1559 goals (informal):

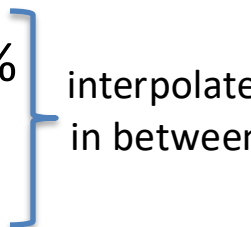
- users incentivized to bid their true utility for posting Tx,
- block proposer incentivized to not create fake Tx, and
- disincentivize off chain agreements

Gas calculation: EIP1559

Every block has a “baseFee”:

the **minimum** gasPrice for all Tx in the block

baseFee is computed from total gas in earlier blocks:

- earlier blocks at gas limit (30M gas) \Rightarrow base fee goes up 12.5%
 - earlier blocks empty \Rightarrow base fee decreases by 12.5%
- 
- interpolate
in between

If earlier blocks at “target size” (15M gas) \Rightarrow base fee does not change

Gas calculation

EIP1559 Tx specifies three parameters:

- **gasLimit**: max total gas allowed for Tx
- **maxFeePerGas**: maximum allowed gas price
- **maxPriorityFeePerGas**: additional “tip” to be paid to block proposer

Computed **gasPrice** bid:

$$\text{gasPrice} = \min(\text{maxFeePerGas}, \text{baseFeePerGas} + \text{maxPriorityFeePerGas})$$

Max Tx fee: **gasLimit** × **gasPrice**

Gas calculation (informal)

gasUsed \leftarrow gas used by Tx

Send **gasUsed** \times (**gasPricePerGas** – **baseFeePerGas**) to proposer

BURN **gasUsed** \times **baseFeePerGas**



\Rightarrow total supply of ETH can decrease

Gas calculation

- (1) if **gasPrice** < **baseFeePerGas**: abort
- (2) If **gasLimit** × **gasPrice** < msg.sender.balance: abort
- (3) deduct **gasLimit** × **gasPrice** from msg.sender.balance

- (4) set **Gas** ← **gasLimit**
- (5) execute Tx: deduct gas from **Gas** for each instruction
if at end (**Gas** < 0): abort, Tx is invalid (proposer keeps **gasLimit** × **gasPrice**)
- (6) Refund **Gas** × **gasPrice** to msg.sender.balance

- (7) **gasUsed** ← **gasLimit** – **Gas**
 - (7a) BURN **gasUsed** × **baseFeePerGas**
 - (7b) Send **gasUsed** × (**gasPricePerGas** – **baseFeePerGas**) to proposer



Base fee
0.7 GWei

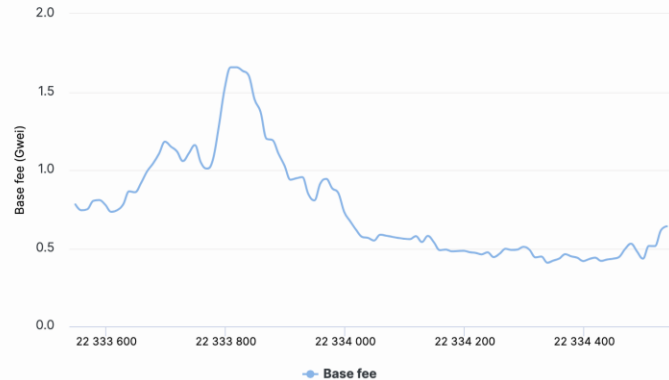
Total burned
5243421.3 ETH
\$9,423,330,940.03

Burn rate (1h)
0.06 ETH/min
\$101.19 / min

Burn rate (24h)
0.14 ETH/min
\$251.09 / min


Gas target
17929741

Block utilization
50.6 %



beaconcha.in

Current ETH Emission



beaconcha.in

Update in: 8s

| Number | Gas Target | Gas Used | Reward | Txs | Age | Base Fee | Burned ETH |
|----------|------------|----------|-------------|-----|-------------|-----------|-------------|
| 22334538 | 17929741 | 10674334 | 0.01686 ETH | 124 | 11 mins ago | 0.68 GWei | 0.00722 ETH |
| 22334537 | 17947267 | 33699919 | 0.01869 ETH | 182 | 12 mins ago | 0.61 GWei | 0.02054 ETH |
| 22334536 | 17964810 | 19121257 | 0.03018 ETH | 225 | 12 mins ago | 0.60 GWei | 0.01156 ETH |
| 22334535 | 17947284 | 1535098 | 0.00195 ETH | 39 | 12 mins ago | 0.68 GWei | 0.00105 ETH |
| 22334534 | 17964827 | 30017637 | 0.03586 ETH | 194 | 12 mins ago | 0.63 GWei | 0.01890 ETH |
| 22334533 | 17982387 | 25937029 | 0.02533 ETH | 160 | 12 mins ago | 0.60 GWei | 0.01548 ETH |
| 22334532 | 17999965 | 12252551 | 0.02630 ETH | 134 | 13 mins ago | 0.62 GWei | 0.00762 ETH |
| 22334531 | 17982405 | 10773799 | 0.00896 ETH | 158 | 13 mins ago | 0.65 GWei | 0.00705 ETH |
| 22334530 | 17999982 | 17913057 | 0.04856 ETH | 162 | 13 mins ago | 0.65 GWei | 0.01173 ETH |

Solidity

<https://docs.soliditylang.org/en/latest/>

Contract structure

```
interface IERC20 {  
    function transfer(address _to, uint256 _value) external returns (bool);  
    function totalSupply() external view returns (uint256);  
    ...  
}  
  
contract ERC20 is IERC20 {      // inheritance  
    address owner;  
    constructor() public { owner = msg.sender; }  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        ... implentation ...  
    }  
}
```

Value types

- uint256
- address (bytes32)
 - `_address.balance`, `_address.send(value)`, `_address.transfer(value)`
 - call: send Tx to another contract

```
bool success = _address.call{value: msg.value/2, gas: 1000}(args);
```
 - delegatecall: load code from another contract into current context
- bytes32
- bool

Reference types

- structs
- arrays
- bytes
- strings
- mappings:

- Declaration: mapping (address => uint256) **balances**;
- Assignment: balances[addr] = value;

```
struct Person {  
    uint128 age;  
    uint128 balance;  
    address addr;  
}  
Person[10] public people;
```

Globally available variables

- **block:** .blockhash, .coinbase, .gaslimit, .number, .timestamp
- gasLeft()
- **msg:** .data, .sender, .sig, .value
- **tx:** .gasprice, .origin
- abi: encode, encodePacked, encodeWithSelector, encodeWithSignature
- Keccak256(), sha256(), sha3()
- **require, assert** e.g.: require(msg.value > 100, "insufficient funds sent")

A → B → C → D:
at D: msg.sender == C
tx.origin == A

Function visibilities

- **external**: function can only be called from outside contract.

Arguments read from calldata

- **public**: function can be called externally and internally.

if called externally: arguments copied from calldata to memory

- **private**: only visible inside contract
- **internal**: only visible in this contract and contracts deriving from it
- **view**: only read storage (no writes to storage)
- **pure**: does not touch storage

```
function f(uint a) private pure returns (uint b) { return a + 1; }
```

Inheritance

```
// SPDX-License-Identifier: MIT
// Compatible with OpenZeppelin Contracts ^5.0.0
pragma solidity ^0.8.27;

import {ERC20} from "@openzeppelin/contracts/token/ERC20/ERC20.sol";
import {Ownable} from "@openzeppelin/contracts/access/Ownable.sol";

contract MyToken is ERC20, Ownable {
    constructor(address initialOwner)
        ERC20("MyToken", "MTK")
        Ownable(initialOwner)
    {}

    function mint(address to, uint256 amount) public onlyOwner {
        _mint(to, amount);
    }
}
```

ERC20 tokens

- <https://github.com/ethereum/EIPs/blob/master/EIPS/eip-20.md>
- A standard API for fungible tokens that provides basic functionality to transfer tokens or allow the tokens to be spent by a third party.
- An ERC20 token is itself a smart contract that maintains all user balances:
mapping(address => uint256) internal **balances**;
- A standard interface allows other contracts to interact with every ERC20 token.
No need for special logic for each token.

ERC20 token interface

- function **transfer**(address _to, uint256 _value) external returns (bool);
- function **transferFrom**(address _from, address _to, uint256 _value) external returns (bool);
- function **approve**(address _spender, uint256 _value) external returns (bool);
- function **totalSupply**() external view returns (uint256);
- function **balanceOf**(address _owner) external view returns (uint256);
- function **allowance**(address _owner, address _spender) external view returns (uint256);

How are ERC20 tokens transferred?

```
contract ERC20 is IERC20 {  
  
    mapping (address => uint256) internal balances;  
  
    function transfer(address _to, uint256 _value) external returns (bool) {  
        require(balances[msg.sender] >= _value, "ERC20_INSUFFICIENT_BALANCE");  
        require(balances[_to] + _value >= balances[_to], "UINT256_OVERFLOW" );  
        balances[msg.sender] -= _value;  
        balances[_to] += _value;  
        emit Transfer(msg.sender, _to, _value);    // write log message  
        return true;  
    }  
}
```

Tokens can be minted by a special function **mint(address _to, uint256 _value)**

ABI encoding and decoding

- Every function has a 4 byte selector that is calculated as the first 4 bytes of the hash of the function signature.
 - For `transfer`, this looks like **`bytes4(keccak256("transfer(address,uint256)"));`**
- The function arguments are then ABI encoded into a single byte array and concatenated with the function selector.
 - This data is then sent to the address of the contract, which is able to decode the arguments and execute the code.
- **Functions can also be implemented within the fallback function**

Calling other contracts

- Addresses can be cast to contract types.

```
address _token;
```

```
IERC20Token tokenContract = IERC20Token(_token);
```

```
ERC20Token tokenContract = ERC20Token(_token);
```

- When calling a function on an external contract, Solidity will automatically handle ABI encoding, copying to memory, and copying return values.
 - **tokenContract.transfer(_to, _value);**

Stack variables

- Stack variables generally cost the least gas
 - can be used for any simple types (anything that is ≤ 32 bytes).
 - `uint256 a = 123;`
- All simple types are represented as `bytes32` at the EVM level.
- Only 16 stack variables can exist within a single scope.

Calldata

- Calldata is a read-only byte array.
- Every byte of a transaction's calldata costs gas
(16 gas per non-zero byte, 4 gas per zero byte).
- It is cheaper to load variables directly from calldata, rather than copying them to memory.
 - This can be accomplished by marking a function as ``external``.

Memory (compiled to MSTORE, MLOAD)

- Memory is a byte array.
- Complex types (anything > 32 bytes such as structs, arrays, and strings) must be stored in memory or in storage.

string memory **name** = "Alice";

- Memory is cheap, but the cost of memory grows quadratically.

Storage array (compiled to SSTORE, SLOAD)

- Using storage is very expensive and should be used sparingly.
- Writing to storage is most expensive.
Reading from storage is cheaper, but still relatively expensive.
- mappings and state variables are always in storage.
- Some gas is refunded when storage is deleted or set to 0
- Trick for saving has: variables < 32 bytes can be packed into 32 byte slots.

Event logs

- Event logs are a cheap way of storing data that does not need to be accessed by any contracts.
- Events are stored in transaction receipts, rather than in storage.

Security considerations

- Are we checking math calculations for overflows and underflows?
 - done by the compiler since Solidity 0.8.
- What assertions should be made about function inputs, return values, and contract state?
- Who is allowed to call each function?
- Are we making any assumptions about the functionality of external contracts that are being called?

Re-entrancy bugs

```
contract Bank{

    mapping(address=>uint) userBalances;

    function getUserBalance(address user) constant public returns(uint) {
        return userBalances[user];    }

    function addToBalance() public payable {
        userBalances[msg.sender] = userBalances[msg.sender] + msg.value;    }

    // user withdraws funds
    function withdrawBalance() public {
        uint amountToWithdraw = userBalances[msg.sender];

        // send funds to caller ... vulnerable!
        if (msg.sender.call{value:amountToWithdraw}() == false) { throw; }
        userBalances[msg.sender] = 0;
    } }
}
```

```

contract Attacker {
    uint numIterations;
    Bank bank;

    function Attacker(address _bankAddress) {    // constructor
        bank = Bank(_bankAddress);
        numIterations = 10;
        if (bank{value:75}.addToBalance() == false)    { throw; }    // Deposit 75 Wei
        if (bank.withdrawBalance() == false)    { throw; }    // Trigger attack
    }

    function () {    // the fallback function
        if (numIterations > 0) {
            numIterations--; // make sure Tx does not run out of gas
            if (bank.withdrawBalance() == false) { throw; }
        }
    }
}

```

Why is this an attack?

(1) Attacker → Bank.addToBalance(75)

(2) Attacker → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance →

Attacker.fallback → Bank.withdrawBalance → ...

withdraw 75 Wei at each recursive step

How to fix?

```
function withdrawBalance() public {  
    uint amountToWithdraw = userBalances[msg.sender];  
  
    userBalances[msg.sender] = 0;  
    if (msg.sender.call{value:amountToWithdraw}() == false) {  
        userBalances[msg.sender] = amountToWithdraw;  
        throw;  
    }  
}
```

END OF LECTURE

Next lecture: DeFi contracts