

Announcement

We're going to use Piazza!

<https://piazza.com/class/m8zibfzpqsi115/>

Find a team by end of week.

Introduction to Bitcoin

Deian Stefan

Slides from Dan Boneh

Today

1. **What is a blockchain?**
2. **Basic introduction to cryptography**
3. **Start looking at Bitcoin (and finish next week)**

What is a blockchain?

user facing tools (cloud servers)

applications (DAPPs, smart contracts)

Execution engine (blockchain computer)

Sequencer: orders transactions

Data Availability / Consensus Layer

Consensus layer (informal)

A public append-only data structure:

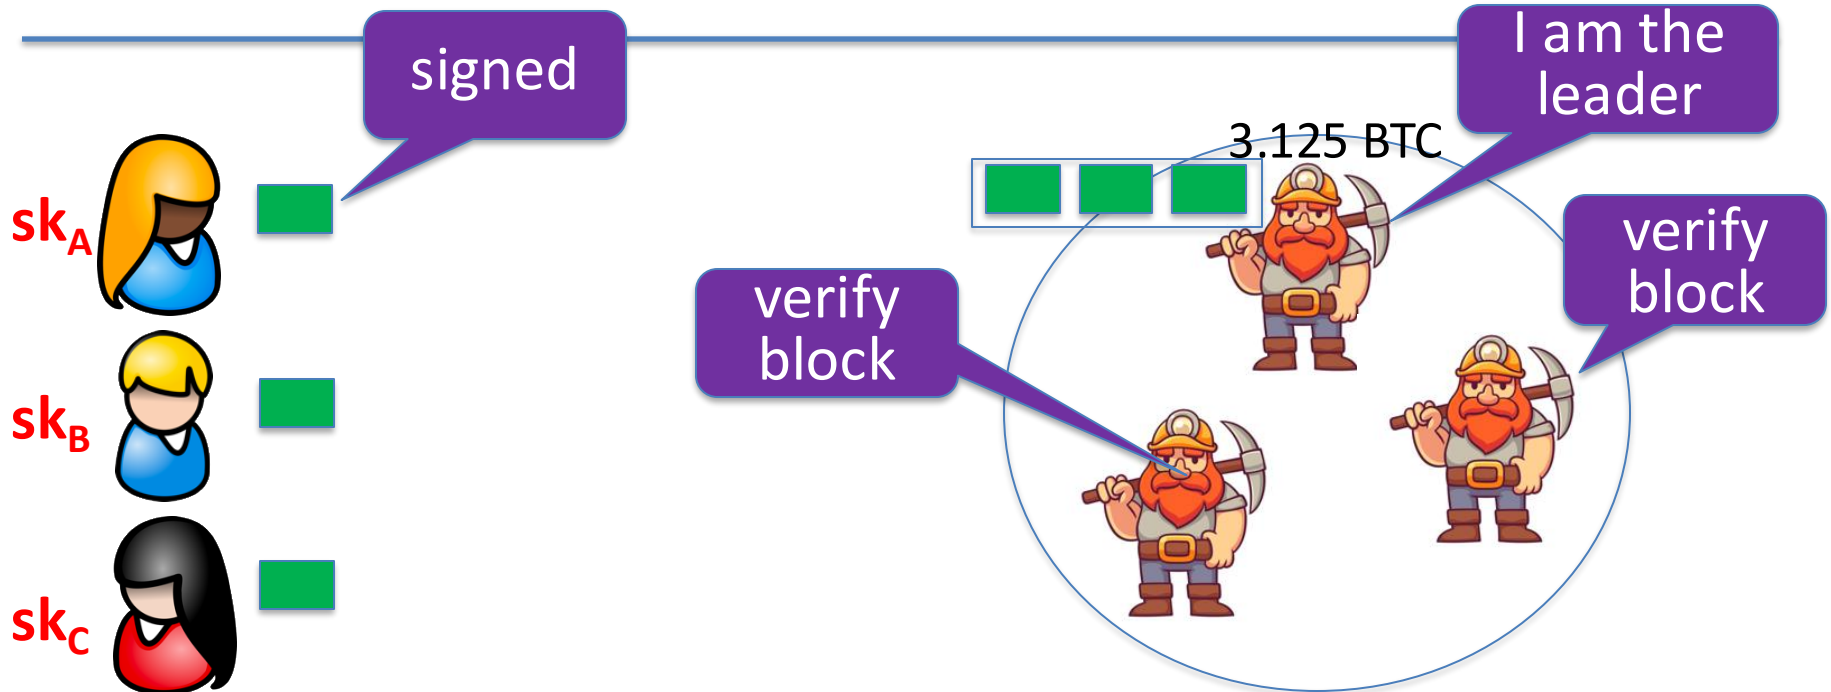
achieved by replication

- **Persistence:** once added, data can never be removed*
- **Safety:** all honest participants have the same data**
- **Liveness:** honest participants can add new transactions
- **Open(?):** anyone can add data (no authentication)

Data Availability / Consensus layer

How are blocks added to chain?

blockchain





How are blocks added to chain?

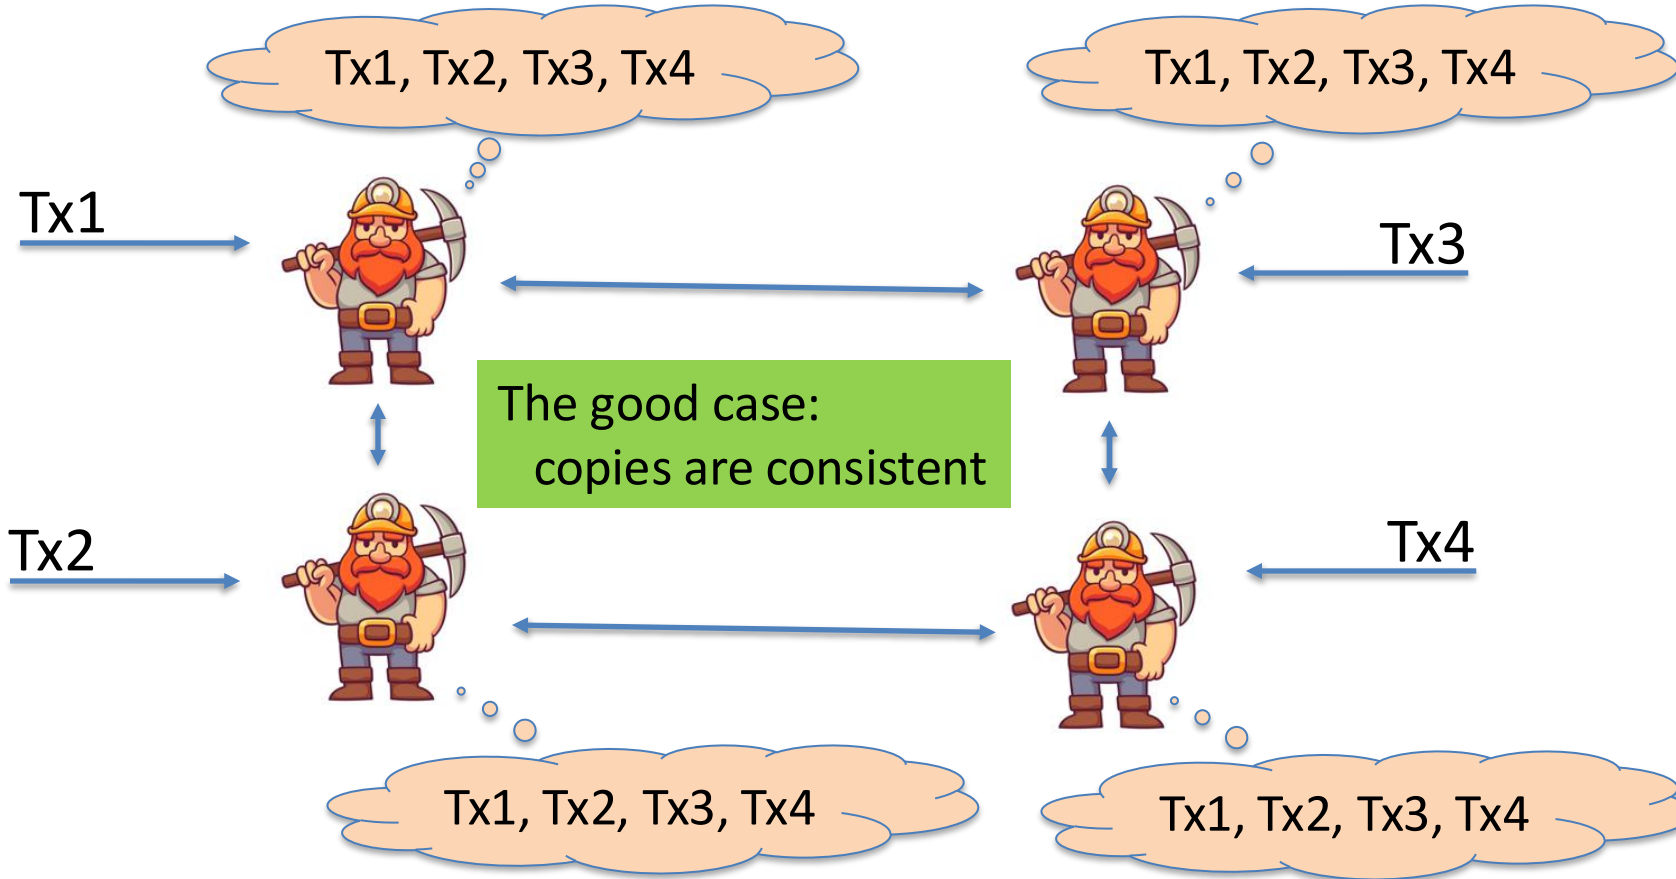
blockchain



...



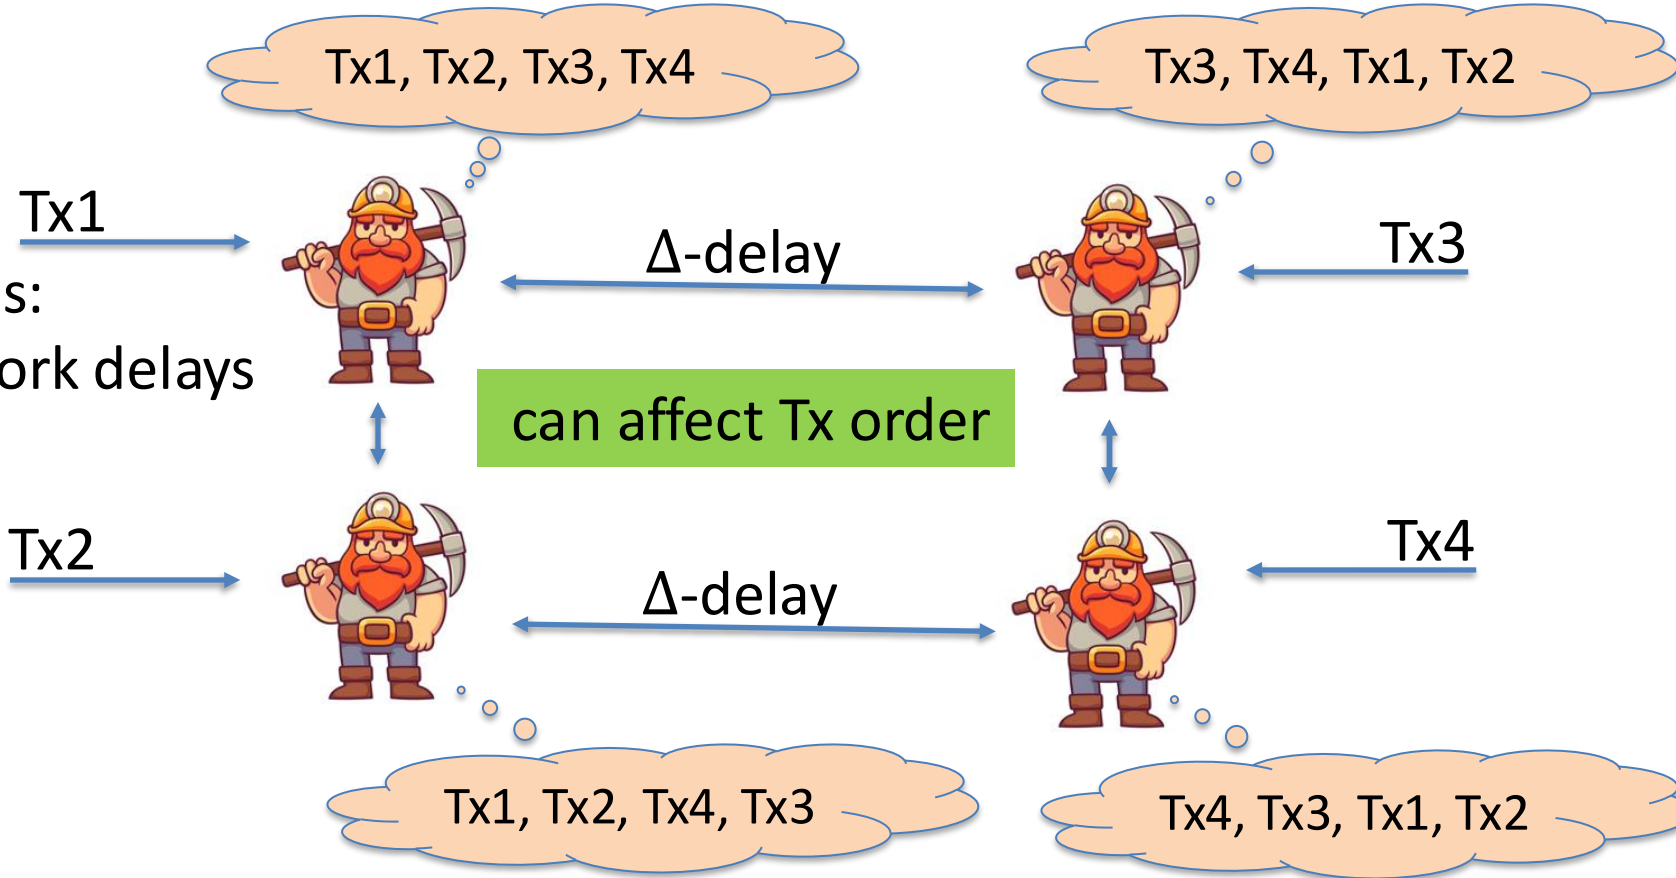
Why is consensus a hard problem?



Why is consensus a hard problem?

Problems:


- Network delays



Why is consensus a hard problem?

Problems:

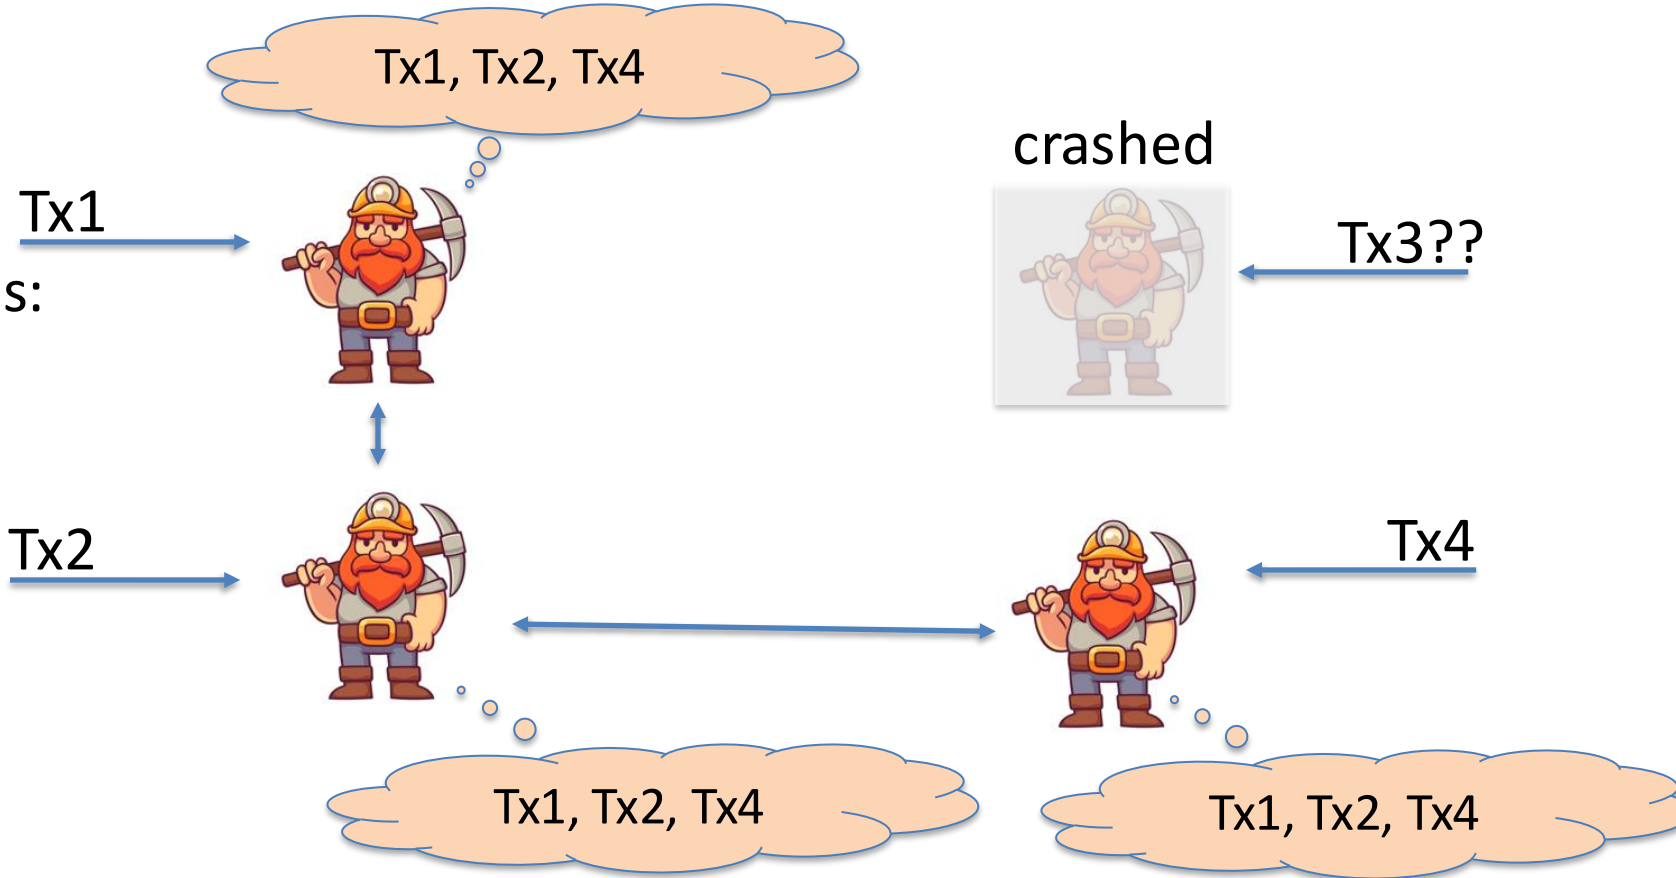
- Network delays
- Network partition



Why is consensus a hard problem?

Problems:

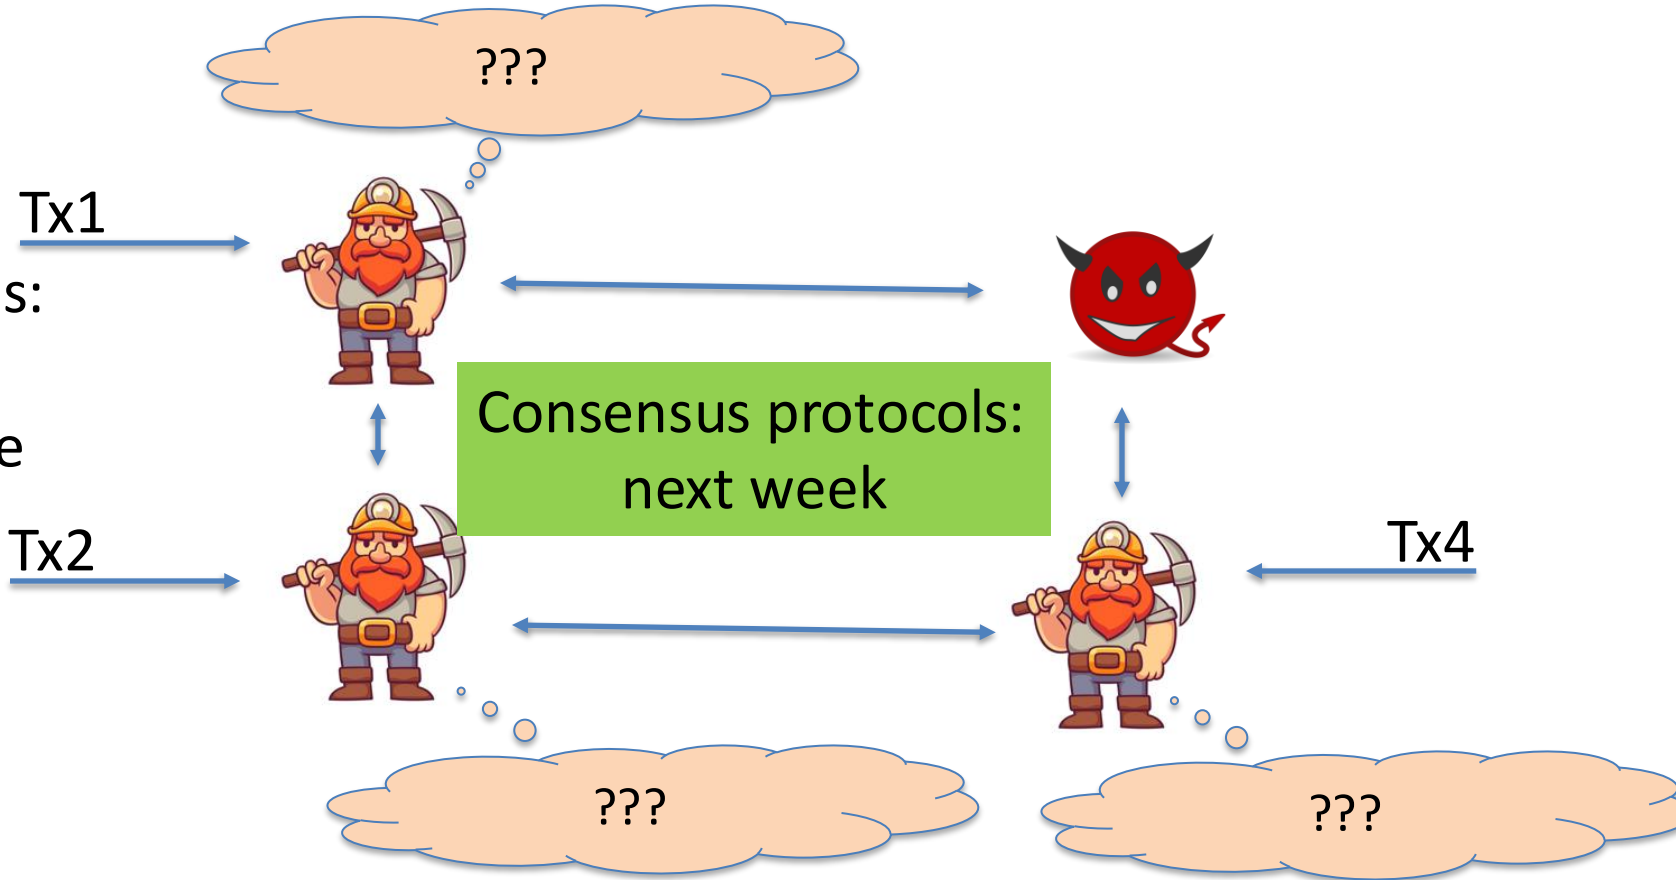
- crash



Why is consensus a hard problem?

Problems:

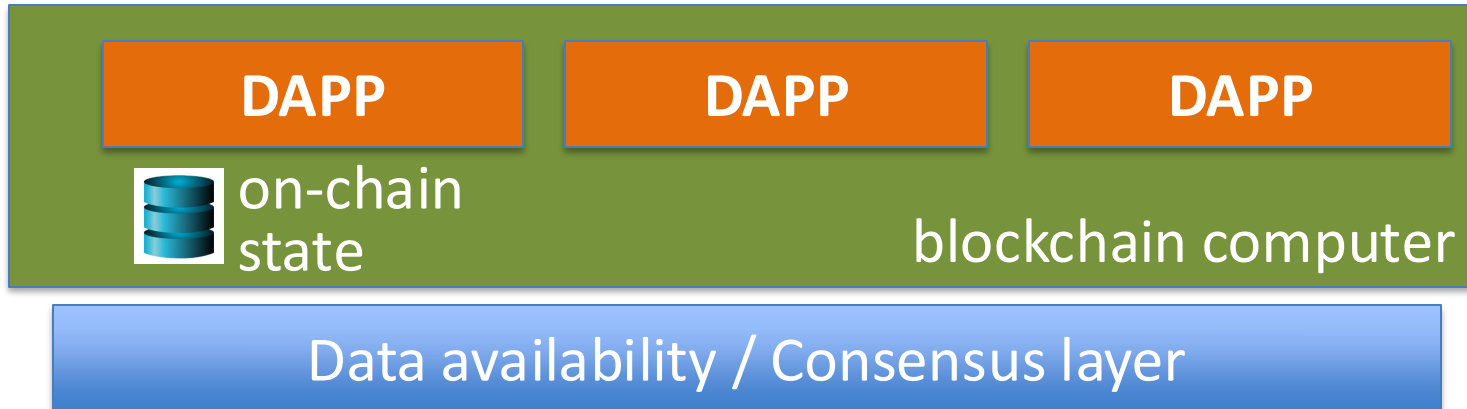
- crash
- malice



Next layer: the blockchain computer

Decentralized applications (DAPPs):

- Run on blockchain: code and state are written on chain
- Accept Tx from users \Rightarrow state transitions are recorded on chain



Next layer: the blockchain computer

Top layer: user facing servers

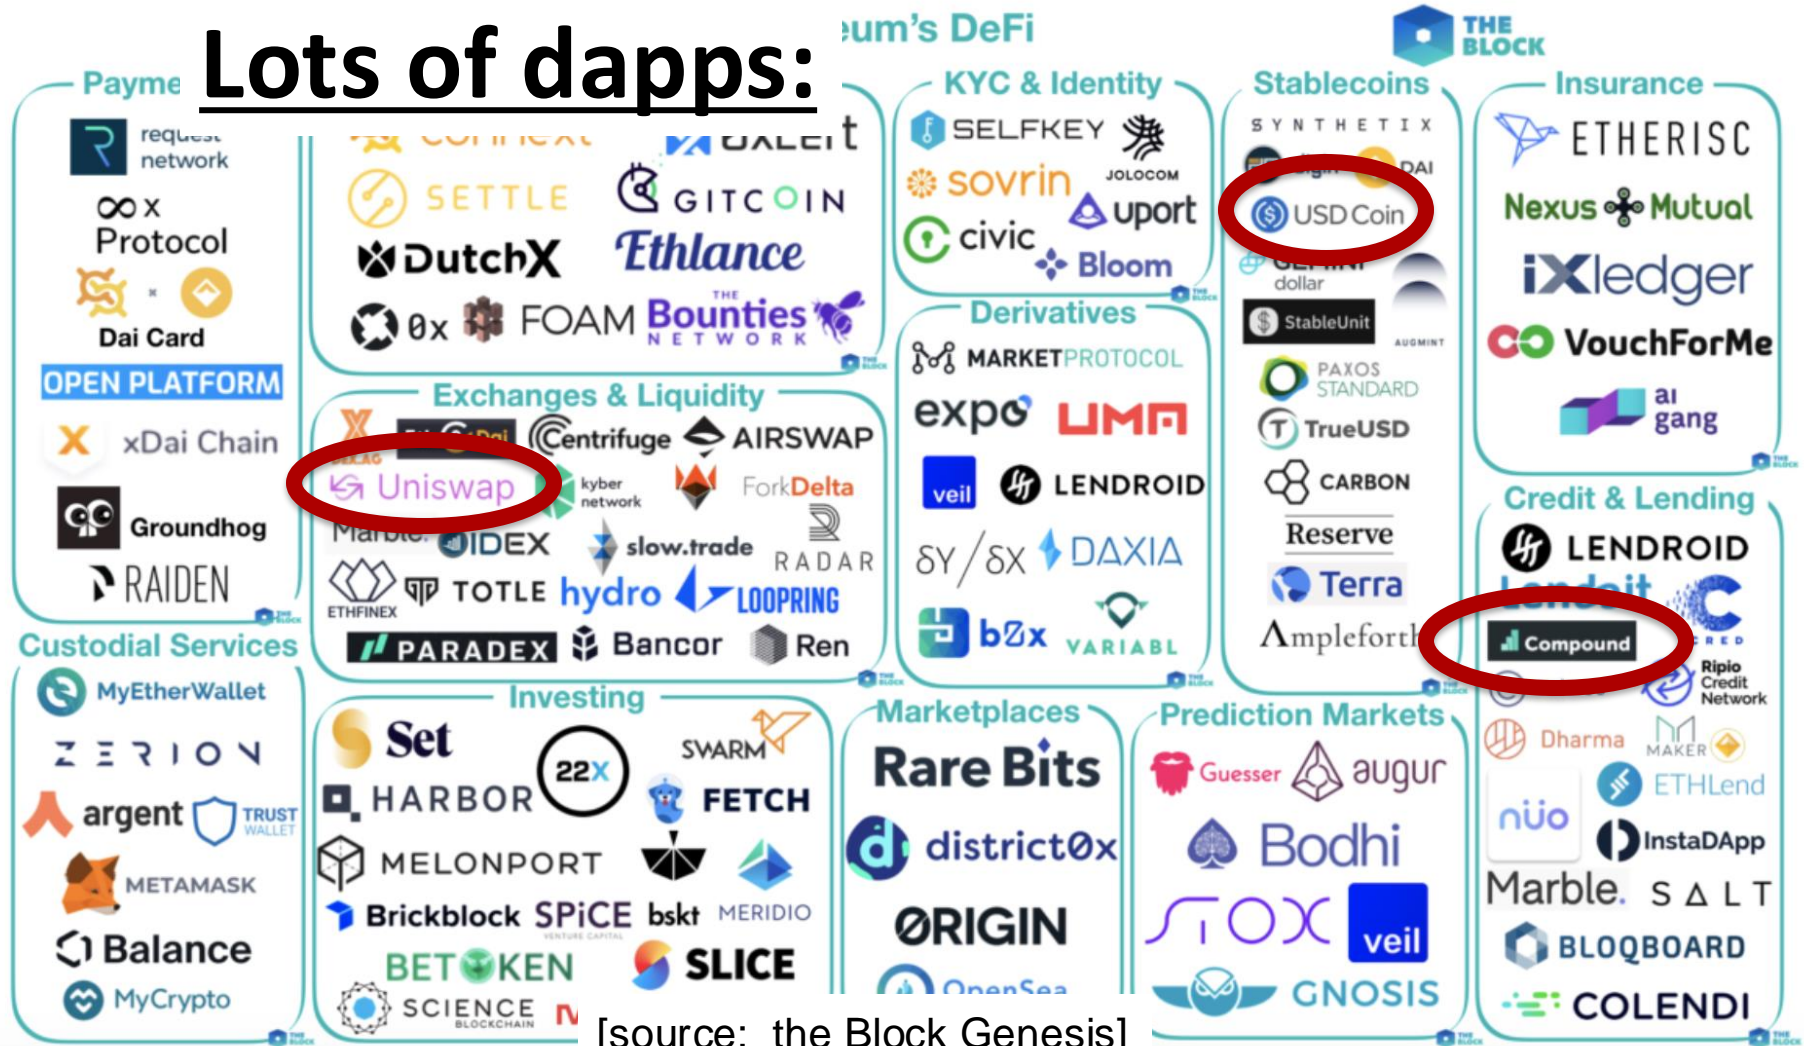


end user



Data availability / Consensus layer

Lots of dapps:




Today

1. What is a blockchain?
2. **Basic introduction to cryptography**
 - Hash functions
 - Commitments
 - Merkle trees
 - Signatures
3. **Start looking at Bitcoin**

Hash functions

An efficiently computable function $H: M \rightarrow T$

where $|M| \gg |T|$



Collision resistance

Def: a collision for $H: M \rightarrow T$ is pair $x \neq y \in M$ s.t. $H(x) = H(y)$

$|M| \gg |T|$ implies that many collisions exist

Def: a function $H: M \rightarrow T$ is collision resistant if it is “hard” to find even a single collision for H (we say H is a CRF)

Example: **SHA256:** $\{x : \text{len}(x) < 2^{64} \text{ bytes}\} \rightarrow \{0,1\}^{256}$

(output is 32 bytes)

Application: committing to data on a blockchain

Alice has a large file m . She posts $h = H(m)$ (32 bytes)

Bob reads h . Later he learns m' s.t. $H(m') = h$

H is a CRF \Rightarrow Bob is convinced that $m' = m$
(otherwise, m and m' are a collision for H)

We say that $h = H(m)$ is a **binding commitment** to m

(note: not hiding, h may leak information about m)

Committing to a list (of transactions)

Alice has $S = (m_1, m_2, \dots, m_n)$

32 bytes



Goal:

- Alice posts a short binding commitment to S , $h = \text{commit}(S)$
- Bob reads h . Given $(m_i, \text{proof } \pi_i)$ can check that $S[i] = m_i$
Bob runs $\text{verify}(h, i, m_i, \pi_i) \rightarrow \text{accept/reject}$

security: adv. cannot find (S, i, m, π) s.t. $m \neq S[i]$ and
 $\text{verify}(h, i, m, \pi) = \text{accept}$ where $h = \text{commit}(S)$

Merkle tree

(Merkle 1989)

commitment



h

Merkle tree
commitment

m_1 m_2 m_3 m_4 m_5 m_6 m_7 m_8



list of values S

Goal:


- commit to list S of size n
- Later prove $S[i] = m_i$

Merkle tree

(Merkle 1989)

[simplified]

commitment



Goal:

- commit to list S of size n
- Later prove $S[i] = m_i$

To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

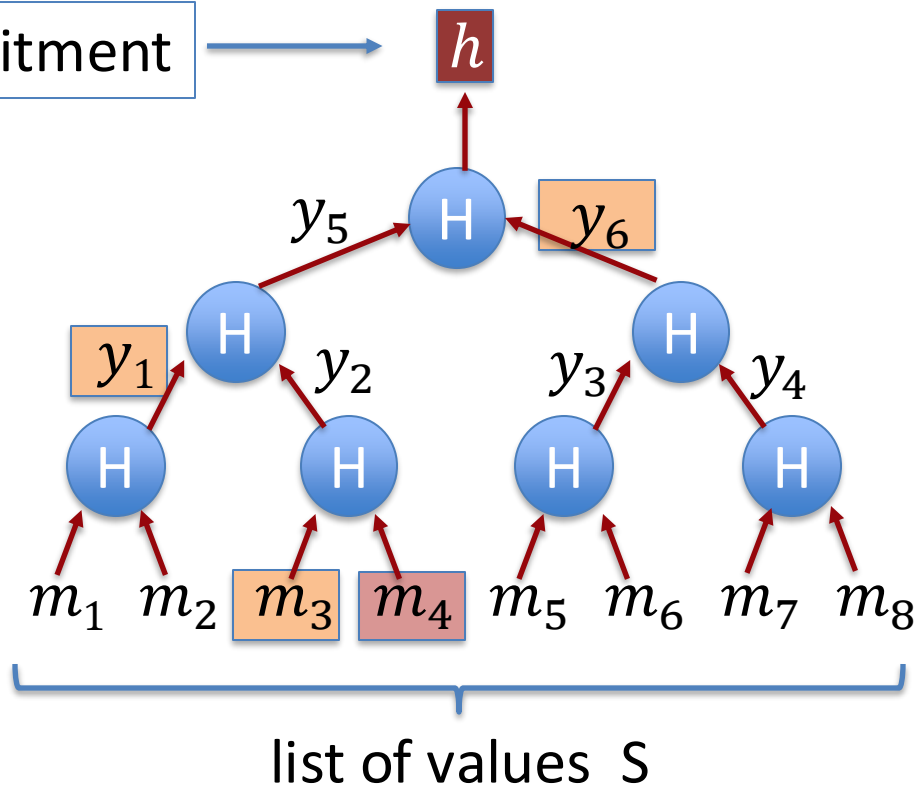
length of proof: $\log_2 n$

Merkle tree

(Merkle 1989)

[simplified]

commitment



To prove $S[4] = m_4$,
proof $\pi = (m_3, y_1, y_6)$

Bob does:

$$y_2 \leftarrow H(m_3, m_4)$$

$$y_5 \leftarrow H(y_1, y_2)$$

$$h' \leftarrow H(y_5, y_6)$$

accept if $h = h'$

Merkle tree

(Merkle 1989)

Thm: For a given n : if H is a CRF then

adv. cannot find (S, i, m, π) s.t. $|S| = n$, $m \neq S[i]$,

$h = \text{commit}(S)$, and $\text{verify}(h, i, m, \pi) = \text{accept}$

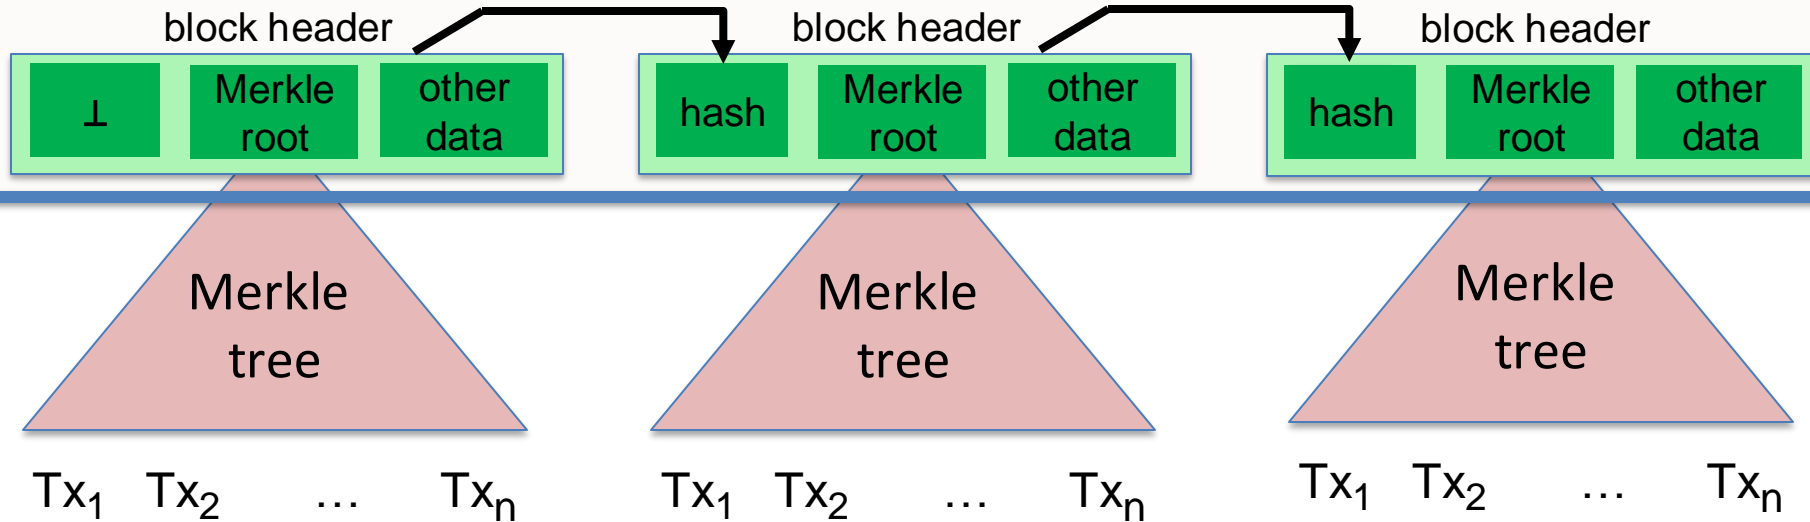
(to prove, prove the contra-positive)

How is this useful? To post a block of transactions S on chain suffices to only write $\text{commit}(S)$ to chain. Keeps chain small.

\Rightarrow Later, can prove contents of every Tx.

Abstract block chain

blockchain



Merkle proofs are used to prove that a Tx is “on the block chain”

Another application: proof of work

Goal: computational problem that

- takes time $\Omega(D)$ to solve, but
- solution takes time $O(1)$ to verify

(D is called the **difficulty**)

How? $H: X \times Y \rightarrow \{0, 1, 2, \dots, 2^n - 1\}$ e.g. $n = 256$

- puzzle: input $x \in X$, output $y \in Y$ s.t. $H(x, y) < 2^n / D$
- verify(x, y): accept if $H(x, y) < 2^n / D$

Another application: proof of work

Thm: if H is a “random function” then the best algorithm requires D evaluations of H in expectation.

Note: this is a parallel algorithm

⇒ the more machines I have, the faster I solve the puzzle.

Proof of work is used in some consensus protocols (e.g., Bitcoin)



Bitcoin uses $H(x, y) = \text{SHA256}(\text{SHA256}(x.y))$

Today

1. What is a blockchain?
2. **Basic introduction to cryptography**
 - Hash functions
 - Commitments
 - Merkle trees
 - **Signatures**
3. **Start looking at Bitcoin**

Signatures

Physical signatures: bind transaction to author

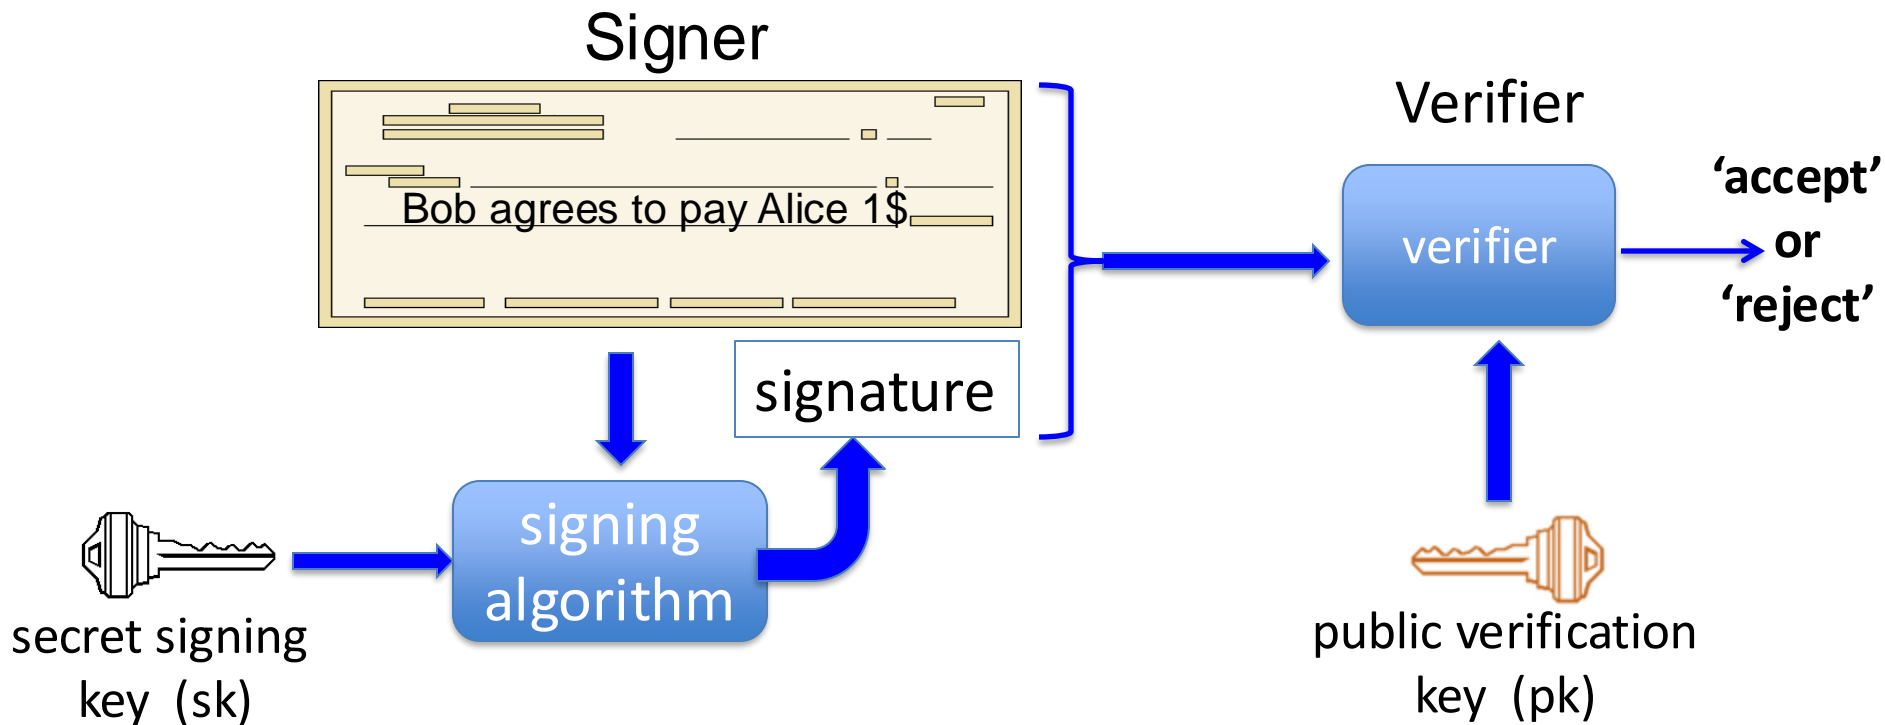


Problem in the digital world:

anyone can copy Bob's signature from one doc to another

Digital signatures

Solution: make signature depend on document



Digital signatures: syntax

Def: a signature scheme is a triple of algorithms:

- **Gen()**: outputs a key pair (pk, sk)
- **Sign**(sk, msg) outputs sig. σ
- **Verify**(pk, msg, σ) outputs 'accept' or 'reject'

Secure signatures: (informal)

Adversary who sees signatures **on many messages** of their choice, cannot forge a signature on a new message.


Families of signature schemes

1. RSA signatures (old ... not used in blockchains):
 - long sigs and public keys (≥ 256 bytes), fast to verify
2. Discrete-log signatures: Schnorr and ECDSA (Bitcoin, Ethereum)
 - short sigs (48 or 64 bytes) and public key (32 bytes)
3. BLS signatures: 48 bytes, aggregatable, easy threshold
(Ethereum 2.0, Chia, Dfinity)

Signatures on the blockchain

Signatures are used everywhere:

- ensure Tx authorization,
- governance votes,
- consensus protocol votes.



In summary ...

Digital signatures: (Gen, Sign, Verify)

$\text{Gen}() \rightarrow (\text{pk}, \text{sk}),$

$\text{Sign}(\text{sk}, m) \rightarrow \sigma, \quad \text{Verify}(\text{pk}, m, \sigma) \rightarrow \text{accept/reject}$

signing key



verification key



Today

1. What is a blockchain?
2. Basic introduction to cryptography
 - Hash functions
 - Commitments
 - Merkle trees
 - Signatures
3. **Start looking at Bitcoin**

Bitcoin mechanics

This lecture: Bitcoin mechanics

user facing tools (cloud servers)

applications (DAPPs, smart contracts)


Execution engine (blockchain computer)

Sequencer: orders transactions


Data Availability / Consensus Layer

start
today

next week



First: overview of the Bitcoin consensus layer




typically, miners are connected to eight other peers (anyone can join)

First: overview of the Bitcoin consensus layer

miners broadcast received Tx
to the P2P network

every miner:
validates received Tx and
stores them in its **mempool**
(unconfirmed Tx)

note: miners see all Tx before
they are posted on chain




First: overview of the Bitcoin consensus layer

blockchain



Every ≈ 10 minutes:

- Each miner creates a candidate block from Tx in its mempool
- a “random” miner is selected (how: next week), and broadcasts its block to P2P network
- all miners validate new block



Bitcoin P2P network


First: overview of the Bitcoin consensus layer

blockchain



Selected miner is paid 3.125 BTC
in **coinbase Tx** (first Tx in the block)

- only way new BTC is created
- block reward halves every four years
⇒ max 21M BTC (currently 19.9M BTC)



note: miner chooses order of Tx in block

Properties (very informal)

Next week:

Safety / Persistence:


- to remove a block, need to convince 51% of mining power *

Liveness:

- to block a Tx from being posted, need to convince 51% of mining power **

(some sub 50% censorship attacks, such as feather forks)

Bitcoin blockchain: a sequence of block headers, 80 bytes each



Bitcoin blockchain: a sequence of block headers, 80 bytes each

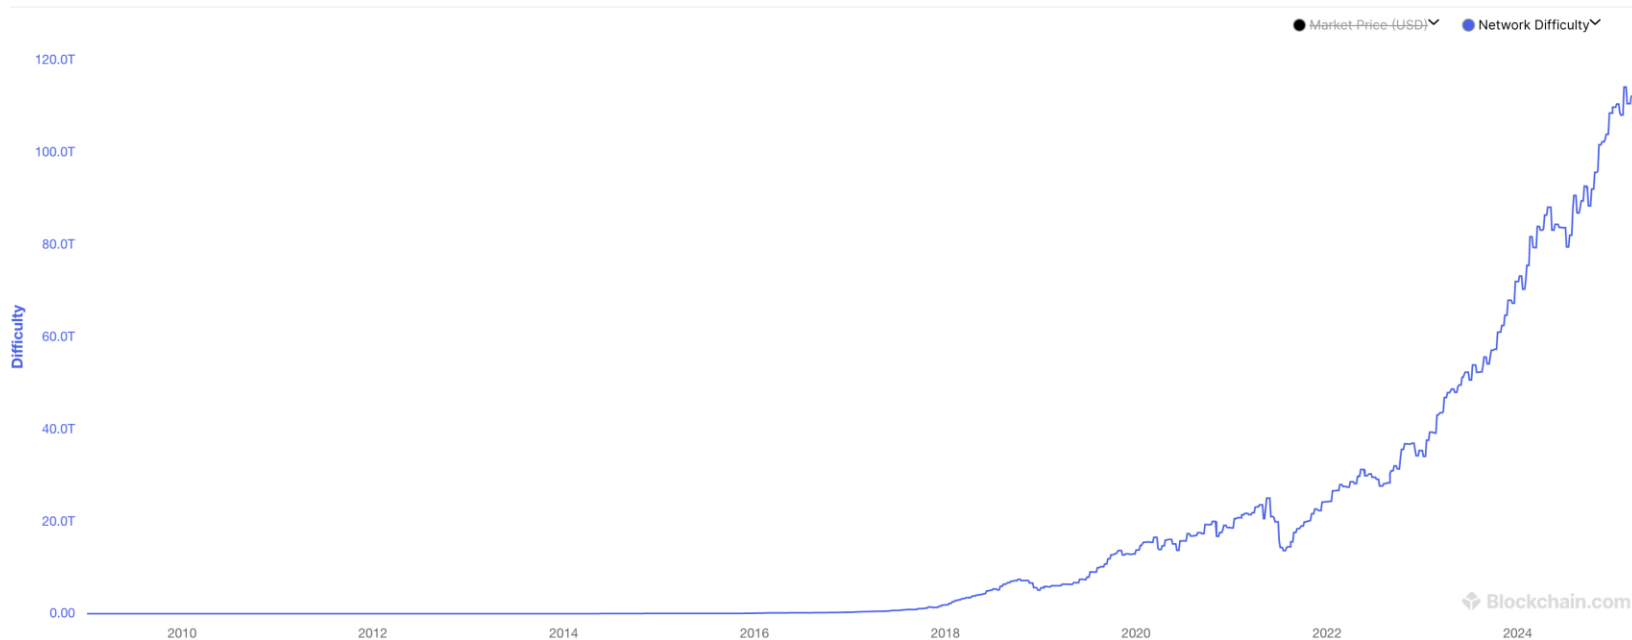
time: time miner assembled the block. Self reported.
(block rejected if too far in past or future)

bits: proof of work difficulty
nonce: proof of work solution } for choosing a leader (next week)

Merkle tree: payer can give a short proof that Tx is in the block

new block every ≈ 10 minutes.


Difficulty over time



An example


Height	Mined	Miner	Size	<div>Tx data </div>	<u>#Tx</u>
648494	17 minutes	Unknown	1,308,663 bytes		1855
648493	20 minutes	SlushPool	1,317,436 bytes		2826
648492	59 minutes	Unknown	1,186,609 bytes		1128
648491	1 hour	Unknown	1,310,554 bytes		2774
648490	1 hour	Unknown	1,145,491 bytes		2075
648489	1 hour	Poolin	1,359,224 bytes		2622

Block 648493


Timestamp	2020-09-15 17:25	
Height	648493	
Miner	SlushPool	(from coinbase Tx)
Number of Transactions	2,826	
Difficulty (D)	17,345,997,805,929.09	(adjusts every two weeks)
Merkle root	350cbb917c918774c93e945b960a2b3ac1c8d448c2e67839223bbcf595baff89	
Transaction Volume	11256.14250596 BTC	
Block Reward		6.25000000 BTC
Fee Reward	this was 2020	0.89047154 BTC (Tx fees given to miner in coinbase Tx)

This lecture


View the blockchain as a sequence of Tx (append-only)



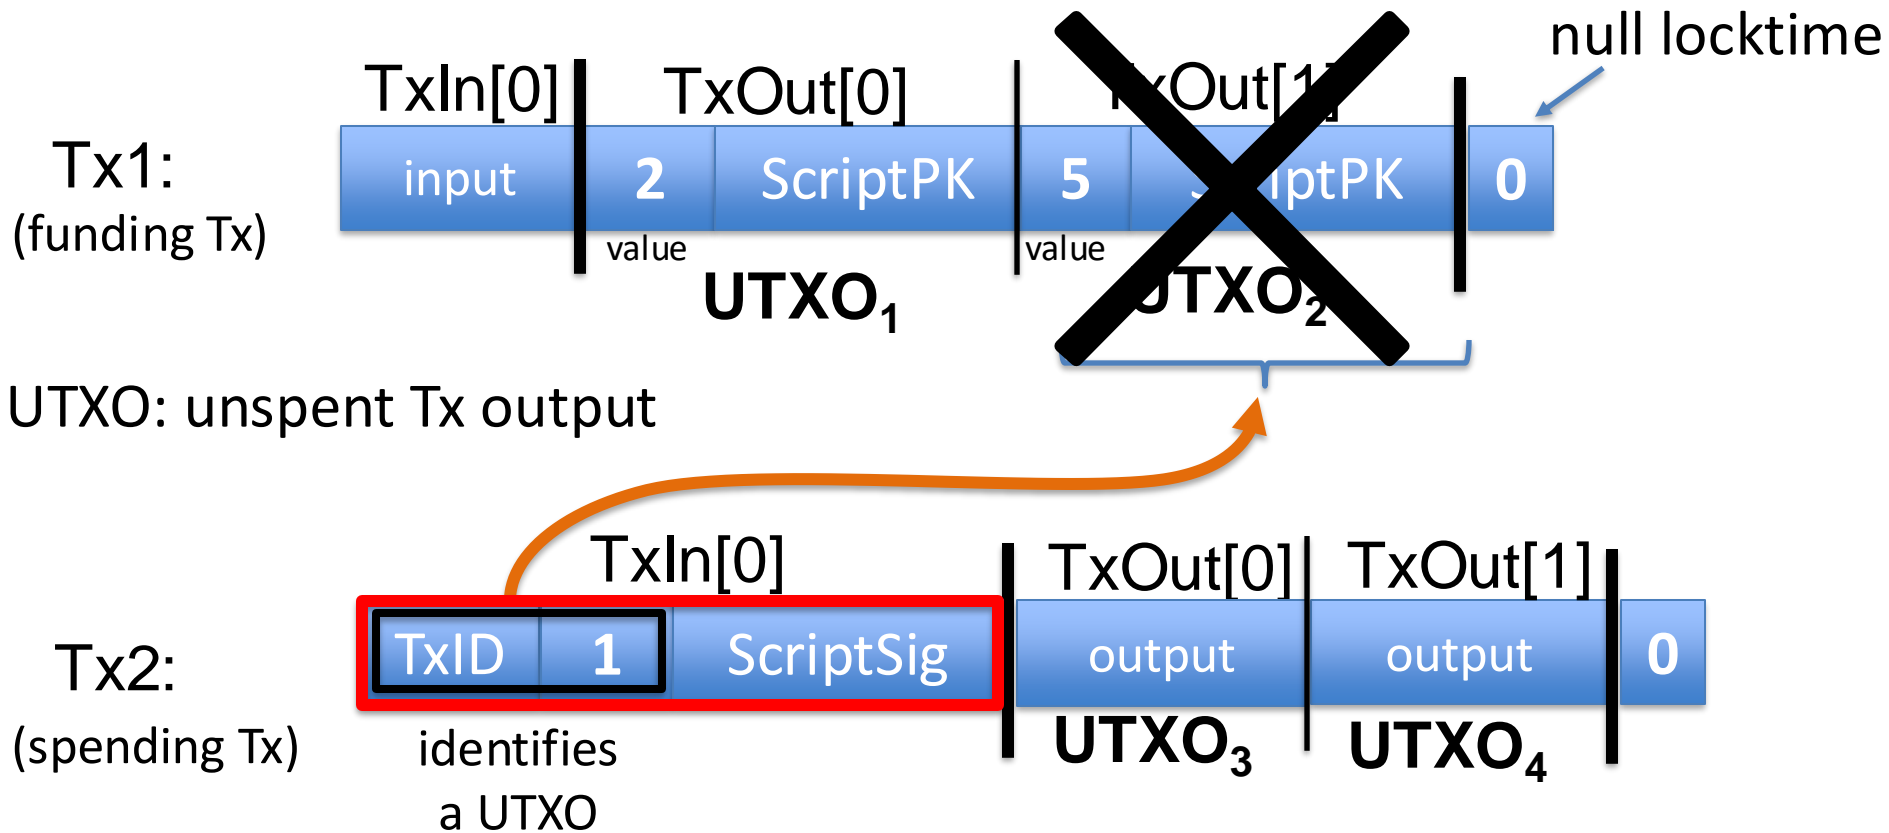
Tx structure (non-coinbase)



Example




Example



Validating Tx2

Miners check (for each input):


program from funding Tx:
under what conditions
can UTXO be spent



1. The program **ScriptSig | ScriptPK** returns true
2. **TxID | index** is in the current UTXO set
3. $\text{sum input values} \geq \text{sum output values}$

After Tx2 is posted, miners remove UTXO_2 from UTXO set

All value in Bitcoin is held in UTXOs



An example (block 648493)

[2826 Tx]

COINBASE (Newly Generated Coins)



1CK6KHY6MHgYvmRQ4PAafKYDrg1ejbH1cE

7.14047154 BTC

OP_RETURN

0.00000000 BTC

OP_RETURN

0.00000000 BTC

Tx0

0.00000000 BTC

6.25 + Tx fees =

7.14047154 BTC

3PuJbxJS1pKxf8EdVR18yBkD1fPAbgUtyw

0.72333974 BTC



1E5Ao1VUnA5BhffvXf2Xmud6avUgwkFnJv

0.00917379 BTC

bc1qr8k3e0vx06lpu3j7m858pa2ak9tyr56ttwvefk 0.61504199 BTC

bc1qdrxve8kua3yz5dgx6wf3u95ngh0d3e648... 0.09290152 BTC

14ZhjuXpQ5jCDjtAy7ZMu3hfEQCWewzLw7 0.00616444 BTC

input

(input UTXO value)

Tx1

outputs

0.72328174 BTC

0.00005800 BTC

(Tx fee)

17MWze4Z1uP1jnvqvj7SAnGtxcoVq11H8A

0.05000000 BTC



3G3C2RFQ8gsf77EQpdR4ZReChWFKEHhxVU

0.04808000 BTC


0.00192000 BTC (Tx fee)

Tx2

0.04808000 BTC

sum of fees in block added to coinbase Tx

Tx fees (all time)



Focusing on Tx2: TxIn[0]

from UTXO
(Bitcoin script)

Value 0.05000000 BTC

Pkscript
OP_DUP
OP_HASH160
45b21c8a0cb687d563342b6c729d31dab58e3a4e
OP_EQUALVERIFY
OP_CHECKSIG

Sigscript
304402205846cace0d73de82dfbdeba4d65b9856d7c1b1730eb401cf4906b2401a69b
dc90220589d36d36be64e774c8796b96c011f29768191abeb7f56ba20ffb0351280860
c01
03557c228b080703d52d72ead1bd93fc72f45c4ddb4c2b7a20c458e2d069c8dd9e

from TxIn[0]

Bitcoin Script

A stack machine. Not Turing Complete: no loops.

Quick survey of op codes:

1. **OP_TRUE** (OP_1), **OP_2**, ..., **OP_16**: push value onto stack

81

82

96

2. **OP_DUP**: push top of stack onto stack

118

Bitcoin Script

3. control:

99 **OP_IF** <statements> **OP_ELSE** <statements> **OP_ENDIF**

105 **OP_VERIFY**: abort fail if top = false

106 **OP_RETURN**: abort and fail

what is this for? ScriptPK = [OP_RETURN, <data>]

136 **OP_EQVERIFY**: pop, pop, abort fail if not equal

Bitcoin Script

4. arithmetic:

OP_ADD, OP_SUB, OP_AND, ...: pop two items, add, push

5. crypto:

OP_SHA256: pop, hash, push

OP_CHECKSIG: pop pk, pop sig, verify sig. on Tx, push 0 or 1

6. Time: **OP_CheckLockTimeVerify** (CLTV):

fail if value at the top of stack > Tx locktime value.

usage: UTXO can specify min-time when it can be spent

Example: a common script

<sig> <pk> **DUP HASH256** <pkhash> **EQVERIFY CHECKSIG**

stack: empty

<sig> <pk>

<sig> <pk> <pk>

<sig> <pk> <hash>

<sig> <pk> <hash> <pkhash>

<sig> <pk>

1

⇒ successful termination

init

push values

DUP

HASH256

push value

EQVERIFY

CHECKSIG


verify(pk, Tx, sig)

Transaction types: (1) P2PKH

pay to public key hash

Alice want to pay Bob 5 BTC:

- step 1: Bob generates sig key pair $(pk_B, sk_B) \leftarrow \text{Gen}()$
- step 2: Bob computes his Bitcoin address as $addr_B \leftarrow H(pk_B)$
- step 3: Bob sends $addr_B$ to Alice
- step 4: Alice posts Tx:




Transaction types: (1) P2PKH

pay to public key hash

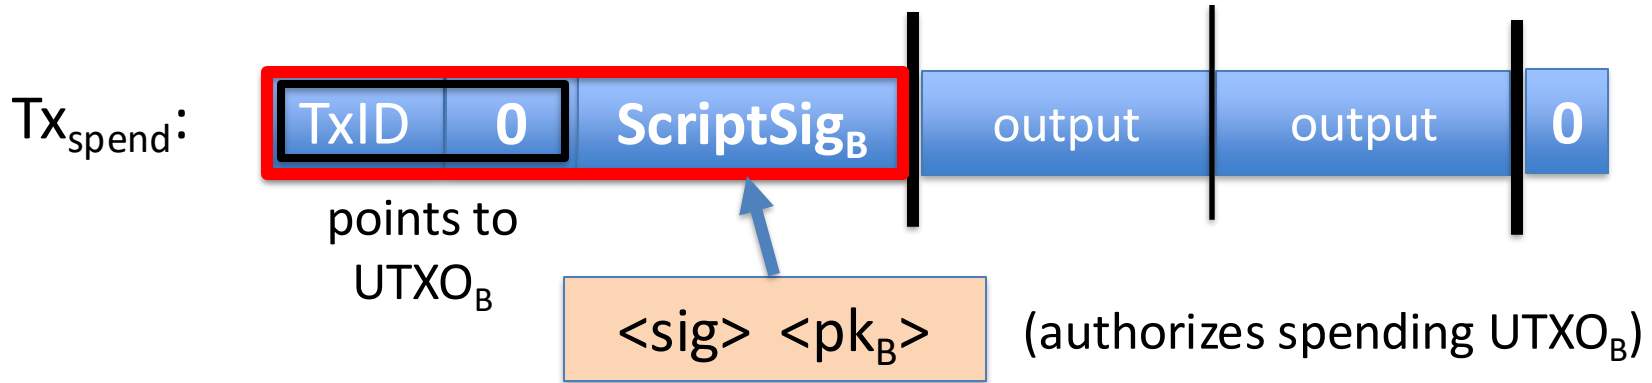
“input” contains ScriptSig that authorizes spending Alice’s UTXO

- example: ScriptSig contains Alice’s signature on Tx
⇒ miners cannot change ScriptPK_B (will invalidate Alice’s signature)



Transaction types: (1) P2PKH

Later, when Bob wants to spend his UTXO: create a Tx_{spend}



$\langle sig \rangle = \text{Sign}(sk_B, Tx)$ where $Tx = (Tx_{\text{spend}} \text{ excluding all ScriptSigs})$ (SIGHASH_ALL)

Miners validate that $ScriptSig_B \mid ScriptPK_B$ returns true

P2PKH: comments

- Alice specifies recipient's pk in UTXO_B
- Recipient's pk is not revealed until UTXO is spent
(some security against attacks on pk)
- Miner cannot change $\langle \text{Addr}_B \rangle$ and steal funds:
invalidates Alice's signature that created UTXO_B

Segregated Witness

ECDSA malleability:

Given (m, sig) anyone can create (m, sig') with $\text{sig} \neq \text{sig}'$

⇒ miner can change sig in Tx and change $\text{TxID} = \text{SHA256}(\text{Tx})$

⇒ Tx issuer cannot tell what TxID is, until Tx is posted

⇒ leads to problems and attacks

Segregated witness: signature is moved to witness field in Tx

$\text{TxID} = \text{Hash}(\text{Tx without witnesses})$

Transaction types: (2) P2SH: pay to script hash

(pre SegWit in 2017)

Let's payer specify a redeem script (instead of just pkhash)

Usage: payee publishes $\text{hash}(\text{redeem script}) \leftarrow \text{Bitcoin addr.}$
payer sends funds to that address

ScriptPK in UTXO: `HASH160 <H(redeem script)> EQUAL`

ScriptSig to spend: `<sig1> <sig2> ... <sign> <redeem script>`

payer can specify complex conditions for when UTXO can be spent

P2SH

Miner verifies:

- (1) $\langle \text{ScriptSig} \rangle \text{ScriptPK} = \text{true}$ \leftarrow payee gave correct script
- (2) $\text{ScriptSig} = \text{true}$ \leftarrow script is satisfied

Example P2SH: multisig

Goal: spending a UTXO requires t-out-of-n signatures

Redeem script for 2-out-of-3: (set by payer)

`<2> <PK1> <PK2> <PK3> <3> CHECKMULTISIG`



hash gives P2SH address

ScriptSig to spend: (by payee)

`<0> <sig1> <sig3> <redeem script>`